

# Solving linear system

**Francesco Sanfilippo**

Istituto Nazionale di Fisica Nucleare - Sezione Roma Tre



27 - 31 March 2017

## Introduction of the problem

- ① Direct solution of linear system  $Ax = b$
- ② Quadratic functional minimization

## Iterative solver

- ① Advantages
- ② Comparison of efficiency

## Checking the solution

- ① Limits, stability and efficiency of various algorithms
- ② Convergence criterions

## Accelerating the convergence

- ① Mixed precision algorithms
- ② Choosing a starting guess
- ③ Preconditioning the problem

## Solving similar problems at the same time

- ① Shifted problems  $A + \sigma \text{Id}$
- ② Deflating the problem

## Review of Parallelisation

- ① Distributed memory
- ② Shared memory
- ③ Vectors

## Gather/scatter approaches

→ 1+2 different examples of gathering of non-local data

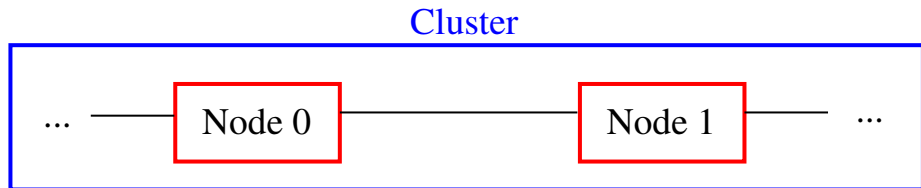
## More specifically on parallelisation

- ① Communication/computation overlap
- ② Multithreading
- ③ Vectorization

## An example of a physical application

→ Lattice QCD

## Lev.1 - Distributed memory paralellism (MPI)



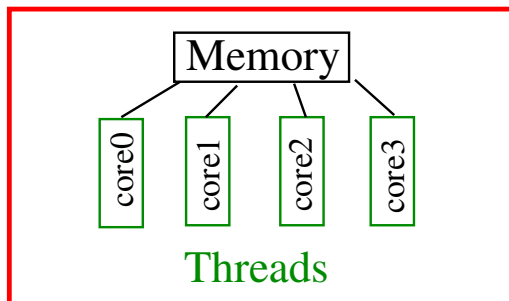
### Features

- Pro:** Using more nodes allow to increase the total computational power
- Con:** Different **Nodes** do not share memory
- Con:** Access to non-local data comes with a **cost** (bandwidth, latency...)

### Tricks

- Arrange computation in such a way to minimize non-local access
- Overlap computation and communication
- Treat different cores inside a cpu with a different approach...

### Single Node



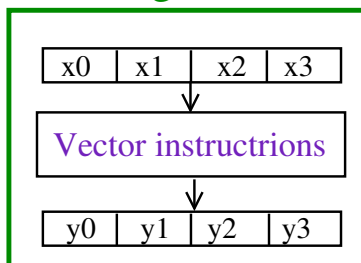
#### Features

- Pro:** Different cores on the same CPU can share the same portion of memory, thus avoiding non-local costly access
- Con:** They can more easily mess-up one with the others trying to access simultaneously to the same data (race condition)

#### Trick

Split the problem in a way that minimise the need for synchronization barriers/mutex

### Single Core



#### Features

**Pro:** Modern cpus include instructions performing many operations at the same time

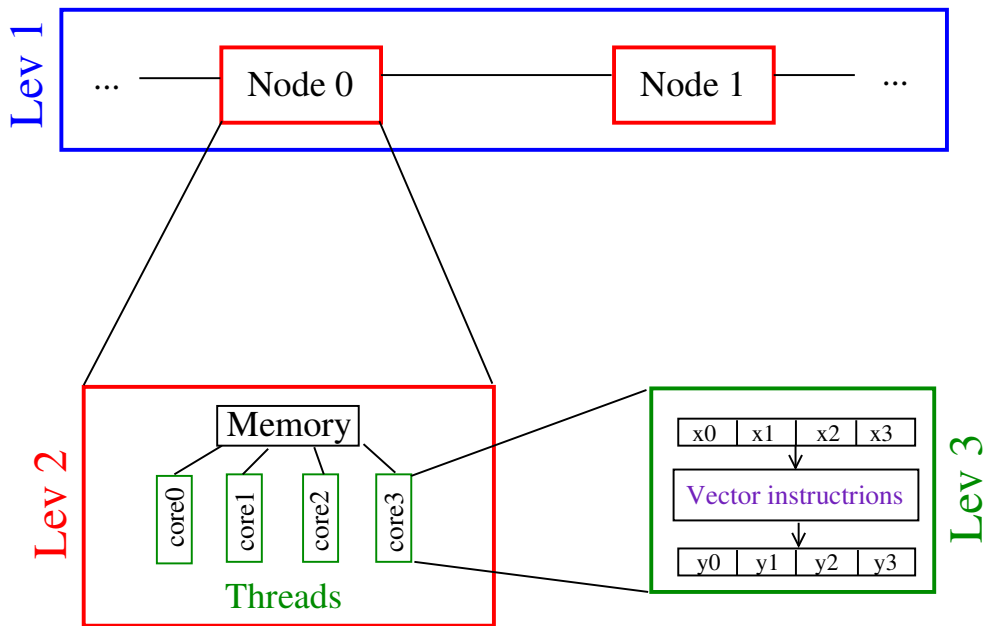
**Con:** Often this restricts to **Same Instruction** being applied to **Many Data** (SIMD)

#### Trick

Layout data in such a way to take advantage of vector instructions



# Three levels of parallelisation



**Target:** Achieve a speed-up as close as possible to the **optimal** one

# 1 dimensional example with 2 nodes

## Distribute vectors through different nodes

Node 0

$p_0$	$p_1$	$p_2$	$p_3$	$p_4$	$\dots$
$= p^0$					

Node 1

$p_{0'}$	$p_{1'}$	$p_{2'}$	$p_{3'}$	$p_{4'}$	$\dots$
$= p^1$					

Vectors are split in the (direct) sum of the sub-vectors distributed to the two nodes  $p = p^0 \oplus p^1$

## Distribute the matrix

$$\begin{bmatrix} t^0 \\ t^1 \end{bmatrix} = \begin{bmatrix} A^{00} & A^{01} \\ A^{10} & A^{11} \end{bmatrix} \times \begin{bmatrix} p^0 \\ p^1 \end{bmatrix}$$

**Problem:** Computing  $b^0$  (on node 0) involves also  $x^1$  (which is in node 1)  
Where to store **different parts** of  $A$ ?

# 1 dimensional example with 2 nodes

## Distribute vectors through different nodes

Node 0

$p_0$	$p_1$	$p_2$	$p_3$	$p_4$	$\dots$
$= p^0$					

Node 1

$p_{0'}$	$p_{1'}$	$p_{2'}$	$p_{3'}$	$p_{4'}$	$\dots$
$= p^1$					

Vectors are split in the (direct) sum of the sub-vectors distributed to the two nodes  $p = p^0 \oplus p^1$

## Distribute the matrix

$$\begin{bmatrix} t^0 \\ t^1 \end{bmatrix} = \begin{bmatrix} A^{00} & A^{01} \\ A^{10} & A^{11} \end{bmatrix} \times \begin{bmatrix} p^0 \\ p^1 \end{bmatrix}$$

**Problem:** Computing  $b^0$  (on node 0) involves also  $x^1$  (which is in node 1)  
Where to store **different parts** of  $A$ ?

$$\begin{bmatrix} t^0 \\ t^1 \end{bmatrix} = \begin{bmatrix} A^{00} & A^{01} \\ A^{10} & A^{11} \end{bmatrix} \times \begin{bmatrix} p^0 \\ p^1 \end{bmatrix}$$

Half of the  $A$  must stored in each node: either divide the rows or the columns

## Computing $t = Ap$

### Local part of the multiplication

- On node 0 we can easily compute  $A^{00}p^0$
- On node 1 we can easily compute  $A^{11}p^1$

## Computing $t = Ap$

### Local part of the multiplication

- On node 0 we can easily compute  $A^{00}p^0$
- On node 1 we can easily compute  $A^{11}p^1$

What about the non-local part?

# Computing $t = Ap$

## Local part of the multiplication

- On node 0 we can easily compute  $A^{00}p^0$
- On node 1 we can easily compute  $A^{11}p^1$

What about the non-local part?

## Local computation of non-local part (gather approach)

- Make each node  $i$  acquire from other node(s)  $j$  the elements needed to compute  $A^{ij}p^j$
- Combine the various parts in place:

$$t^i = \underbrace{A^{ii}t^i}_{\text{local}} + A^{ij} \underbrace{p^j}_{\text{gathered}}$$

# Computing $t = Ap$

## Local part of the multiplication

- On node 0 we can easily compute  $A^{00}p^0$
- On node 1 we can easily compute  $A^{11}p^1$

What about the non-local part?

## Local computation of non-local part (gather approach)

- Make each node  $i$  acquire from other node(s)  $j$  the elements needed to compute  $A^{ij}p^j$
- Combine the various parts in place:

$$t^i = \underbrace{A^{ii}t^i}_{\text{local}} + A^{ij} \underbrace{p^j}_{\text{gathered}}$$

## Non-Local computation of non-local part (scatter approach)

- Make each node  $j$  compute  $A^{ij}p^j$  for all other node(s)  $i$
- Combine the various parts after receiving them:

$$t^i = \underbrace{A^{ii}p^i}_{\text{local}} + \underbrace{A^{ij}p^j}_{\text{scattered}}$$

Not much of difference if the matrix  $A$  is dense

# Computing $t = Ap$

## Local part of the multiplication

- On node 0 we can easily compute  $A^{00}p^0$
- On node 1 we can easily compute  $A^{11}p^1$

What about the non-local part?

## Local computation of non-local part (gather approach)

- Make each node  $i$  acquire from other node(s)  $j$  the elements needed to compute  $A^{ij}p^j$
- Combine the various parts in place:

$$t^i = \underbrace{A^{ii}t^i}_{\text{local}} + A^{ij} \underbrace{p^j}_{\text{gathered}}$$

## Non-Local computation of non-local part (scatter approach)

- Make each node  $j$  compute  $A^{ij}p^j$  for all other node(s)  $i$
- Combine the various parts after receiving them:

$$t^i = \underbrace{A^{ii}p^i}_{\text{local}} + \underbrace{A^{ij}p^j}_{\text{scattered}}$$

Not much of difference if the matrix  $A$  is dense

What if  $A$  is sparse?



## Sparse matrix - Laplace equation problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & s \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

## Sparse matrix - Laplace equation problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} \textcolor{red}{d} & \textcolor{red}{s} & 0 & | & 0 & 0 & s \\ \textcolor{red}{s} & \textcolor{red}{d} & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} \textcolor{red}{p}_0^0 \\ \textcolor{red}{p}_1^0 \\ \textcolor{red}{p}_2^0 \\ \hline p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

$A^{00} p^0$

This part of the multiplication involves only local data

## Sparse matrix - Laplace equation problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & \textcolor{green}{s} \\ s & d & s & | & 0 & 0 & 0 \\ \hline 0 & \textcolor{green}{s} & \textcolor{green}{d} & | & \textcolor{green}{s} & 0 & 0 \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline \textcolor{green}{p}_0^1 \\ p_1^1 \\ \textcolor{green}{p}_2^1 \end{pmatrix}$$

$A^{00}p^0$

This part of the multiplication involves only local data

$A^{01}p^1$

- This part of the multiplication involves non-local data
- Not all the entries of  $p^1$  contributes to  $A^{01}p^1$ !

## Sparse matrix - Laplace equation problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & \textcolor{green}{s} \\ s & d & s & | & 0 & 0 & 0 \\ \hline 0 & \textcolor{green}{s} & d & | & \textcolor{green}{s} & 0 & 0 \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline \textcolor{green}{p}_0^1 \\ p_1^1 \\ \textcolor{green}{p}_2^1 \end{pmatrix}$$

$A^{00}p^0$

This part of the multiplication involves only local data

$A^{01}p^1$

- This part of the multiplication involves non-local data
- Not all the entries of  $p^1$  contributes to  $A^{01}p^1$ !

Only  $sp_0^1$  and  $sp_2^1$  must be computed

**Gather inputs:** collect  $p_0^1$  and  $p_2^1$  on node 0 and multiply them by  $s$

or

**Scatter outputs:** compute  $sp_0^1$  and  $sp_2^1$  on node 1 and send it to node 0

The two approaches are equally good in this case

## Sparse matrix - A slightly different problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & \underline{s} & \underline{s} \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{s} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ s & 0 & s & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline \underline{p_0^1} \\ \underline{p_1^1} \\ \underline{p_2^1} \end{pmatrix}$$

$A^{00}p^0$

This part of the multiplication involves only local data

$A^{01}p^1$

- This part of the multiplication involves non-local data
- Now all the entries of  $p^1$  contributes to  $A^{01}p^1$ !

## Sparse matrix - A slightly different problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & \underline{s} & \underline{s} \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{s} & \underline{0} \\ \hline 0 & 0 & s & | & d & s & 0 \\ s & 0 & s & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline \underline{p_0^1} \\ \underline{p_1^1} \\ \underline{p_2^1} \end{pmatrix}$$

$A^{00}p^0$

This part of the multiplication involves only local data

$A^{01}p^1$

- This part of the multiplication involves non-local data
- Now all the entries of  $p^1$  contributes to  $A^{01}p^1$ !

$sp_0^1$ ,  $sp_1^1$  and  $sp_2^1$  must be computed

**Gather inputs:** collect  $p_0^1$ ,  $p_1^1$  and  $p_2^1$  on node 0, combine them:  $sp_0^1 + sp_1^1$  and  $sp_1^1 + sp_2^1$

or

**Scatter outputs:** compute  $sp_0^1 + sp_1^1$  and  $sp_1^1 + sp_2^1$  on node 1 and send it to node 0

## Sparse matrix - A slightly different problem

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & \underline{s} & \underline{s} \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{s} & \underline{0} \\ \hline 0 & 0 & s & | & d & s & 0 \\ s & 0 & s & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline \underline{p_0^1} \\ \underline{p_1^1} \\ \underline{p_2^1} \end{pmatrix}$$

$A^{00}p^0$

This part of the multiplication involves only local data

$A^{01}p^1$

- This part of the multiplication involves non-local data
- Now all the entries of  $p^1$  contributes to  $A^{01}p^1$ !

$sp_0^1$ ,  $sp_1^1$  and  $sp_2^1$  must be computed

**Gather inputs:** collect  $p_0^1$ ,  $p_1^1$  and  $p_2^1$  on node 0, combine them:  $sp_0^1 + sp_1^1$  and  $sp_1^1 + sp_2^1$

or

**Scatter outputs:** compute  $sp_0^1 + sp_1^1$  and  $sp_1^1 + sp_2^1$  on node 1 and send it to node 0

The second approach involves less communications (the inverse could happen)

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & s \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

Let's try to keep the approach "as simple as possible"

## Minimal modification to the scalar code

- Whenever a component is needed, ask for it to the appropriate node
- When computing  $t$  components requiring non-local  $p$ , write to the node holding it and ask it to send this piece of info



$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & s \\ s & d & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

Let's try to keep the approach "as simple as possible"

## Minimal modification to the scalar code

- Whenever a component is needed, ask for it to the appropriate node
- When computing  $t$  components requiring non-local  $p$ , write to the node holding it and ask it to send this piece of info

## Communication problems

- All nodes have to periodically listen for request
- Problem: **deadlocks**

## Gather Version 0 - Avoiding deadlocks

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} d & s & 0 & | & 0 & 0 & s \\ s & d & s & | & 0 & 0 & 0 \\ \hline \underline{0} & \underline{s} & \underline{d} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & d & s & 0 \\ 0 & 0 & 0 & | & s & d & s \\ s & 0 & 0 & | & 0 & s & d \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

Let's try to keep the approach "as simple as possible"

### Deadlock

If two nodes are simultaneously asking to each other some non-local data they get **stuck!**

### Solution

(Valid for THIS operator that communicate only with first neighbors)

- ① Fix parity of nodes (even/odd)
- ② Split the application of the matrix in two steps:
  - ① in the first, EVEN nodes listen while ODD ones ask
  - ② in the second they switch of role



## Efficiency

- Lot of overhead
- Many branches
- 1/2 nodes are always idle

## Scalability

- Difficult to thread
- Difficult to generalize to arbitrary number of nodes

# Gather Approach 1 - Shift the vector to apply

## Diagonal part

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} \mathbf{d} & s & 0 & | & 0 & 0 & s \\ s & \mathbf{d} & s & | & 0 & 0 & 0 \\ \underline{0} & \underline{s} & \underline{\mathbf{d}} & | & \underline{s} & \underline{0} & \underline{0} \\ 0 & 0 & s & | & \mathbf{d} & s & 0 \\ 0 & 0 & 0 & | & s & \mathbf{d} & s \\ s & 0 & 0 & | & 0 & s & \mathbf{d} \end{pmatrix} \begin{pmatrix} p_0^0 \\ p_1^0 \\ p_2^0 \\ \hline p_0^1 \\ p_1^1 \\ p_2^1 \end{pmatrix}$$

Start applying **diagonal part** (local by definition)

**SHIFT** the vector to apply:  $vp' = p_{i+1}$

$$\begin{pmatrix} t_0^0 \\ t_1^0 \\ t_2^0 \\ \hline t_0^1 \\ t_1^1 \\ t_2^1 \end{pmatrix} = \begin{pmatrix} \mathbf{s} & 0 & 0 & | & 0 & s & d \\ d & \mathbf{s} & 0 & | & 0 & 0 & s \\ \underline{s} & \underline{d} & \underline{\mathbf{s}} & | & \underline{0} & \underline{0} & \underline{0} \\ 0 & s & d & | & \mathbf{s} & 0 & 0 \\ 0 & 0 & s & | & d & \mathbf{s} & 0 \\ 0 & 0 & 0 & | & s & d & \mathbf{s} \end{pmatrix} \begin{pmatrix} p_1^0 \\ p_2^0 \\ p_0^1 \\ \hline p_1^1 \\ p_2^1 \\ p_0^0 \end{pmatrix}$$

Entries above the matrix are now on the diagonal (therefore all operations are local)

## Gather Approach 1 - Shift the vector to apply

### Shifting backward (to shift forward do the opposite)

- Create a buffer vector  $p'$
- Copy local data:  $p'_i = p_{i+1}$  if  $i + 1$  is local
- Communicate backward the first local element and receive the last one from forward node

## Gather Approach 1 - Shift the vector to apply

### Shifting backward (to shift forward do the opposite)

- Create a buffer vector  $p'$
- Copy local data:  $p'_i = p_{i+1}$  if  $i + 1$  is local
- Communicate backward the first local element and receive the last one from forward node

### Avoiding buffer and local copy

- Introduce a shifted local coordinate:  $s(i) = (i + \delta) \bmod N_{loc}$
- Every time we shift backward:  $\delta = \delta + 1$
- Every time we shift forward:  $\delta = \delta - 1$
- We leave local data where it is but we replace
  - $p_{s(N_{loc}-1)}^n$  with  $p_{s(0)}^{n+1}$  when we shift backward
  - $p_{s(0)}^n$  with  $p_{s(N_{loc}-1)}^{n-1}$  when we shift forward

# Gather Approach 1 - Shift the vector to apply

## Pro

- Easy to program, no deadlock
- No branch
- No idle node
- Easy to generalize to different operator (do more shifts)
- Communications are all done in a single bunch for each shift (good for regular patterns)
- Consecutive access to memory (only if really shifting the vector)

## Con

Requires:

- To take a copy of the source vector (more memory) or
- To undo the shift to come back to original vector (more communications)



## Gather Approach 2 - Caching non-local data

$$(p_0^0 \ p_1^0 \ p_2^0)$$

### Initialization

- 1 Identify all the needed component before applying the matrix
- 2 Send to each node in advance the list of component that he will have to send
- 3 Set a buffer for non-local data, keep track of things to receive from/send to other nodes

### At each application

- 1 Collect all data needed from other nodes
- 2 Simultaneously send to other nodes what they need in turn
- 3 Apply the matrix

### Where to store non-local data?

- We can store non-local data together with local one
- Define the cached buffer at the end/origin of the local vector
- We then can easily point to cache when non-local data is required (and no branch occur!)

$$(0 \mid p_0^0 \ p_1^0 \ p_2^0 \mid 0) \xrightarrow{comm} (p_2^1 \mid p_0^0 \ p_1^0 \ p_2^0 \mid p_0^1)$$

## Gather Approach 2 - Caching non-local data

### Pro

- 1 single block of communication
- no copy involved

### Con

- Harder to program (when in multiple dimensions and using cartesian grid)
- Needs more book-keeping (for more complicated operators)
- Cached data can take significant amount of memory
- Read from memory in a fragmented way

Very few modifications are needed

## At the beginning

- Vectors must be distributed
- Matrix must be distributed or its parameter broadcasted

## At each iteration

- Global scalar products must be computed
- Application of the matrix to the vector must be performed according to one of the proposed approach